

CSCI 544

Applied Natural Language Processing

Mohammad Rostami
USC Computer Science Department



Logistical Notes

- **HW1::**
 - Report: explaining how you solve the problem along with the requested output + Jupyter Notebook with printed output + executable .py (different software versions are OK but specify the version clearly in your report)
- **Project Group Formation Deadline: 09/12**
 - We will form random assignment after this date (30 groups so far)
 - Check slack/Piazza and the Excel sheet (Group of 3 or more on Excel)
 - Meet weekly, helpful for both HW and project
- **Paper Selection Deadline: 09/19**
 - Focus on project topic
 - Last year projects: YouTube
 - Read papers from venues such as EMNLP, ACL, NAACL, etc. (often needed in industry as well)
 - Pick your paper soon:
https://docs.google.com/spreadsheets/d/1_vafG77ijmETCnuVZvKpT35k--5op5wn71GZXgAY7O0


Linear models

- A linear function in n -dimensional space (i.e. we have n features) is defined by $n+1$ weights:

$$Y = \sum_{i=0}^n \beta_i X_i$$


- We find the model weights such that the linear function acts as a good predictive model

- Is not necessarily unique!




Perceptron

- Invented by Frank Rosenblatt in 1969
- Inspired by the nervous system
- Unit-based: analogous to a neural cell
- Model: neural activity is modeled by mathematical operations



Perceptron


- First neural network learning model in the 1960's
- Simple and limited (single layer models)
- Basic concepts are similar for multi-layer models so this is a good learning tool
- Still used in current applications (modems, etc.)
- We use discretization to solve classification tasks



$$z = \begin{cases} 1 & \text{if } \sum_{i=1}^n x_i w_i \geq \theta \\ 0 & \text{if } \sum_{i=1}^n x_i w_i < \theta \end{cases}$$

Perceptron Learning Algorithm

- An iterative algorithm:
 - We initialize weights with random values
 - We do several pass on the whole training dataset one by one and update the weights
 - We will not necessarily find a unique solution
- Least perturbation principle
 - Only change weights if there is an error
 - Use small *learning rate* (η) to change weights sufficiently to make the current pattern correct
- Each iteration through the training set is an *epoch*



○ Class 0
● Class 1

Perceptron Learning Algorithm

- Let w_i be the weights vector at iteration, l be a constant (the learning rate), y be the label for the current iteration, z be the current **thresholded** output, and x be the input

$$\beta_{i+1} = \beta_i + \Delta\beta$$
$$\Delta\beta = l(y - z)x$$

y	z	$y - z$
0	0	0
0	1	-1
1	0	1
1	1	0


- Continue training until total training set error ceases to improve
- Perceptron Convergence Theorem: Guaranteed to find a solution in finite time if a solution exists

Perceptron Learning Example

We would like a perceptron to correctly classify the five 2-dimensional data points below.

Let the random initial weight vector $\beta_0 = (1, -2, 2)^T$, why? ($ax+by+c = 0$ or $[a \ b \ c][x \ y \ 1]^T = 0$)

Then the dividing line crosses at $(0, 1)^T$ and $(-2, 0)^T$.



○ Class 0

● Class 1

Let us pick the misclassified point $(-2, -1)^T$ for learning assuming $l = 1$:

$\mathbf{x}_1 = [-2, -1, 1]^T$ (include offset 1)

$\mathbf{z}_1 = [1, -2, 2] \cdot (-2, -1, 1)^T = 2 \rightarrow 1$ (\mathbf{x}_1 is in class 0) ($y_1 - z_1 = -1$)

$\beta_1 = \beta_0 - \mathbf{x}_1$ (let us set $l = 1$ for simplicity)

$\beta_1 = (1, -2, 2)^T - (-2, -1, 1)^T = (3, -1, 1)^T$

The new dividing line crosses at $(0, 1)^T$ and $(-1/3, 0)^T$.

Perceptron Learning Example

Let us pick the misclassified point $(-1, 2)^T$ for learning:


$$\mathbf{x}_2 = [-1, 2, 1]^T \text{ (include offset 1)}$$

$$\mathbf{z}_2 = [3, -1, 1] \cdot (-1, 2, 1)^T = -4 \rightarrow 0 \text{ (}\mathbf{x}_2 \text{ is in class 1) (}\mathbf{y}_2 - \mathbf{z}_2 = 1\text{)}$$

$$\boldsymbol{\beta}_2 = \boldsymbol{\beta}_1 + \mathbf{x}_2 \text{ (let us set } l = 1 \text{ for simplicity)}$$

$$\boldsymbol{\beta}_2 = (3, -1, 1)^T + (-1, 2, 1)^T = (2, 1, 2)^T$$

The new dividing line crosses at $(0, -2)^T$ and $(-1, 0)^T$.




○ Class 0

● Class 1

Perceptron Learning Example


- The algorithm has converged in one epoch because all the points are classified correctly
- In general, we may need several epochs
- Learning rate value is important for fast convergence
- At each epoch, we may use a different random order on the data points
- Neural networks, including the state-of-the-art deep networks are extensions of perceptron (potentially using different learning algorithms)



○ Class 0
● Class 1

Optimal Boundary

- The perceptron learning algorithm is guaranteed to find a solution if the data points are linearly separable
- But are perceptron solutions optimal?




Perfect classification of training samples but may not generalize well to new (untrained) samples.

- Idea: compute a continuous, differentiable error function between input and desired output
 - Define an objective function to measure quality of a model
 - Find the weights for which the objective function is minimal.
 - With a differential function, we can use the gradient descent techniques to find a good solution (at least locally optimal)

Gradient decent

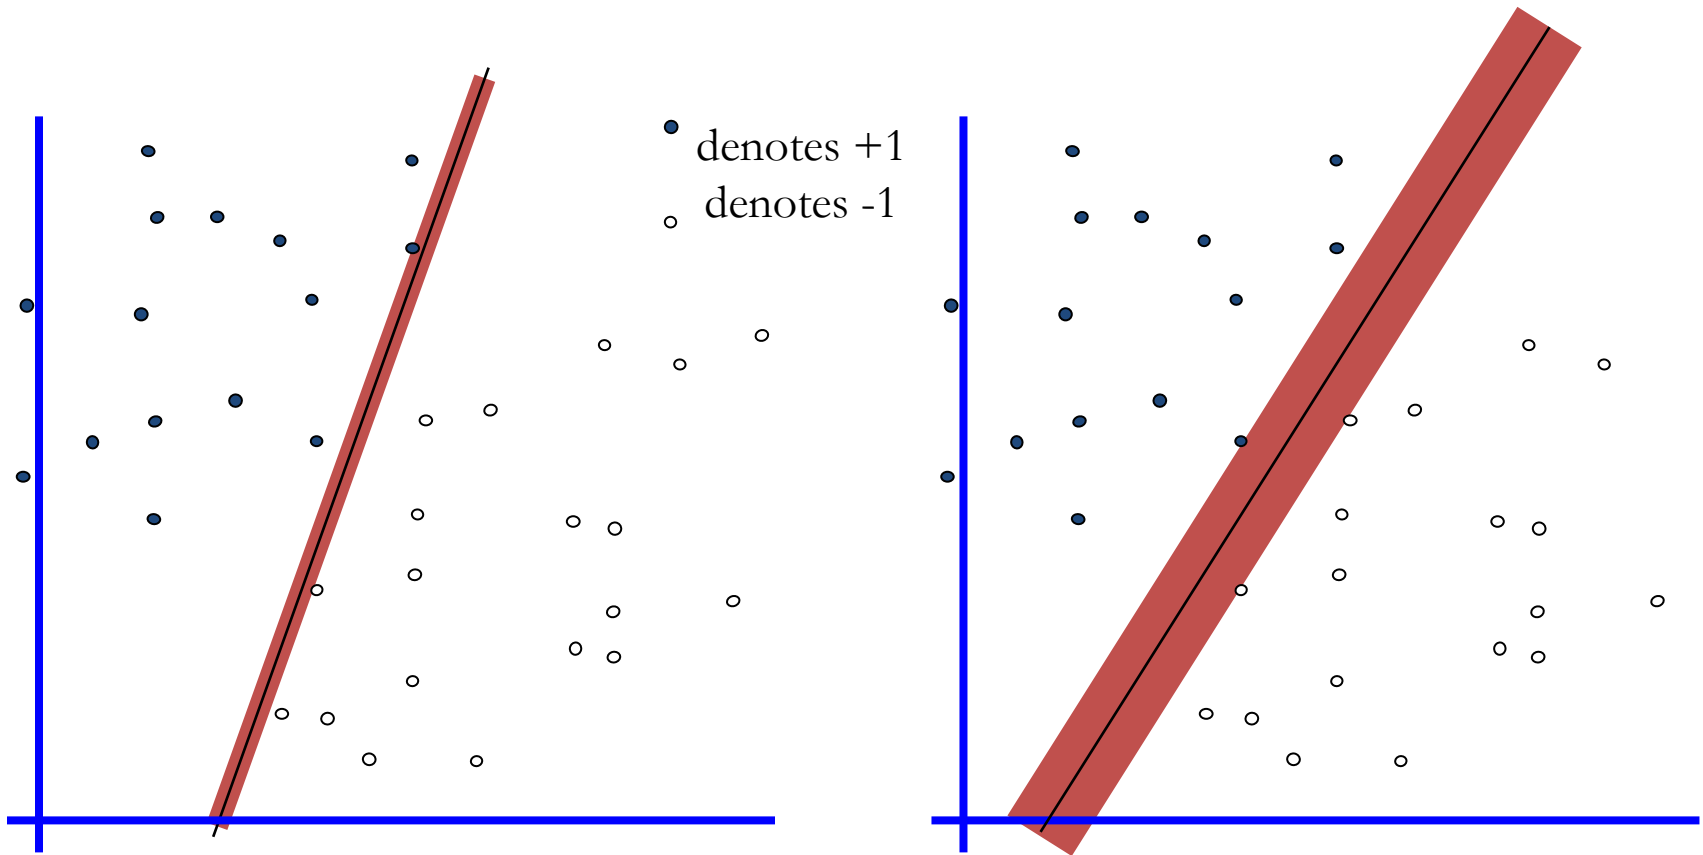
- Suitable for finding minimums of convex differentiable functions (cost functions)



$$\theta_j = \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$$

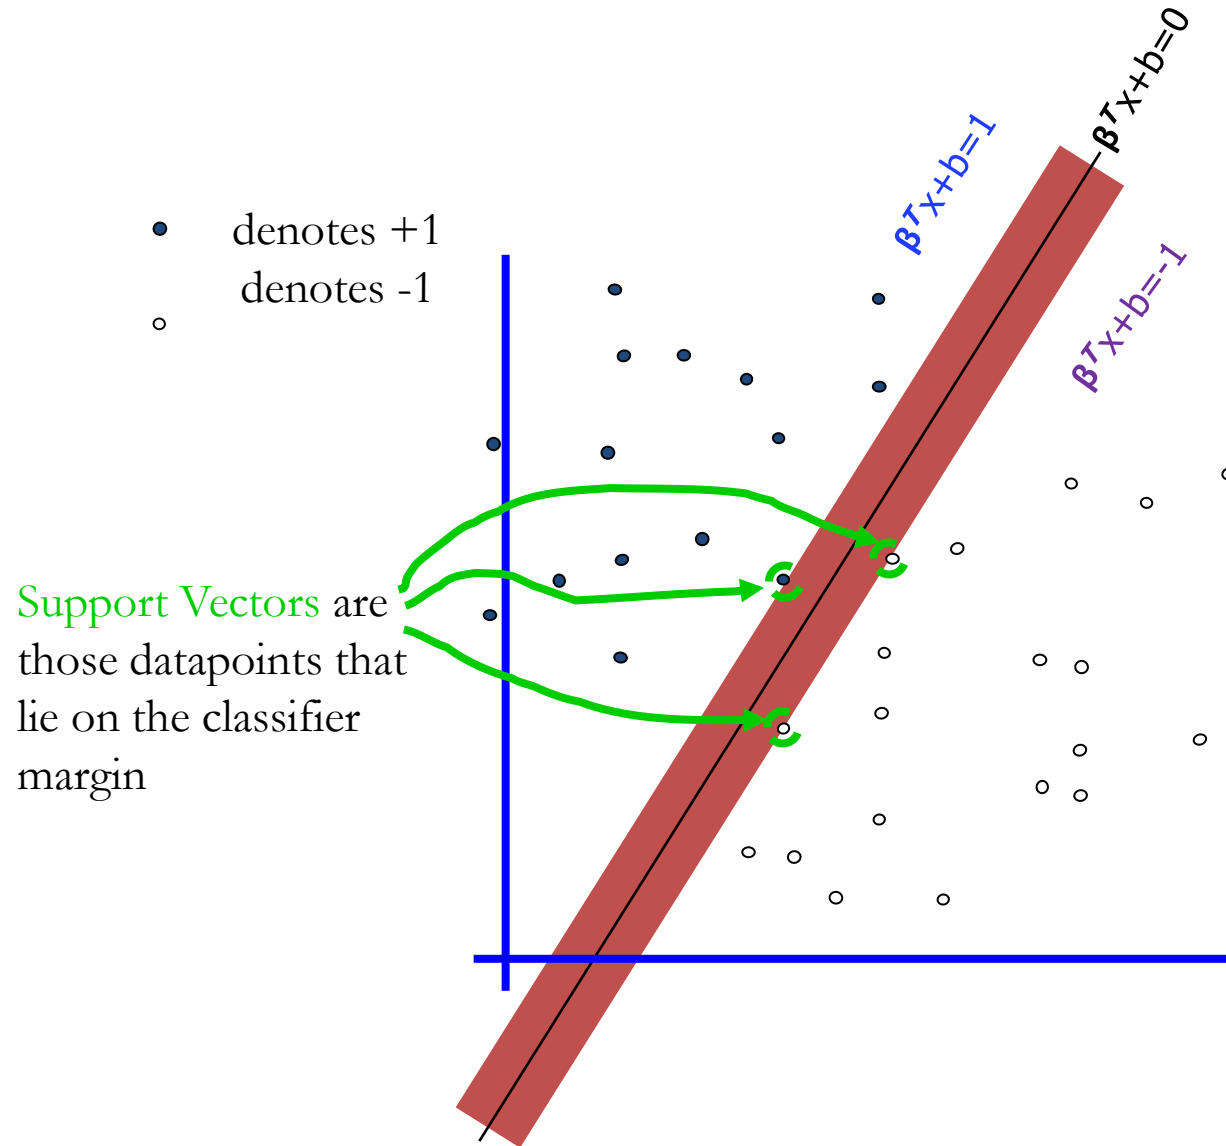
Support Vector Machine

The **margin** of a linear classifier as the width that the boundary could be increased by before hitting a datapoint.



Support Vector Machine (SVM): the **maximum margin linear classifier**

Support Vector Machine




Computing the Margin Width

- Plus-plane = $\{ \mathbf{x} : \boldsymbol{\beta}^T \cdot \mathbf{x} + b = +1 \}$ $M = | \mathbf{x}^+ - \mathbf{x}^- | = | \lambda \boldsymbol{\beta} |$
- Minus-plane = $\{ \mathbf{x} : \boldsymbol{\beta}^T \cdot \mathbf{x} + b = -1 \}$
- Let \mathbf{x}^- be any point on the minus plane
- Let \mathbf{x}^+ be the closest plus-plane-point to \mathbf{x}^- .
- We can deduce:

$$\begin{array}{l} \boldsymbol{\beta}^T \cdot \mathbf{x}^+ + b = +1 \\ \boldsymbol{\beta}^T \cdot \mathbf{x}^- + b = -1 \end{array} \longrightarrow \boldsymbol{\beta}^T \cdot (\mathbf{x}^+ - \mathbf{x}^-) = 2 \longrightarrow \boldsymbol{\beta}^T \cdot (\lambda \boldsymbol{\beta}) = 2 \longrightarrow \lambda = \frac{2}{\boldsymbol{\beta}^T \cdot \boldsymbol{\beta}}$$

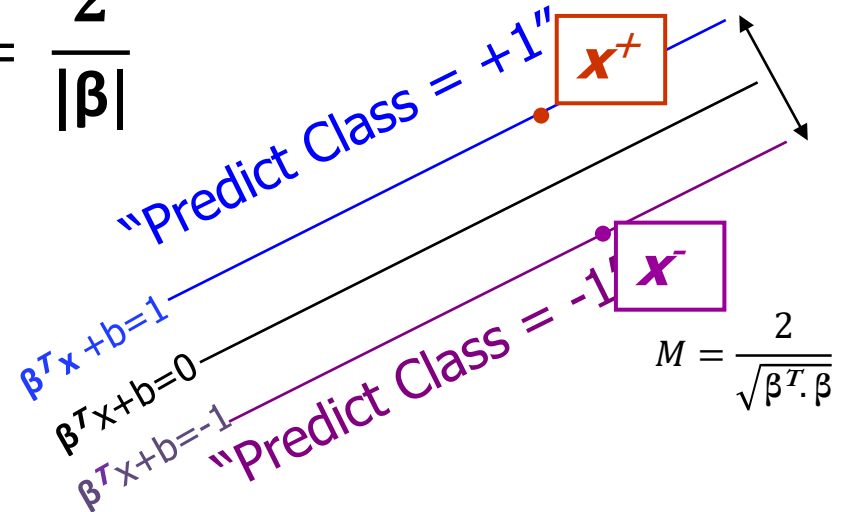
$$M = | \mathbf{x}^+ - \mathbf{x}^- |$$

$$\mathbf{x}^+ = \mathbf{x}^- + \lambda \boldsymbol{\beta}$$



Computing the margin width

$$M = |\mathbf{x}^+ - \mathbf{x}^-| = |\lambda \boldsymbol{\beta}| = \left| \frac{2}{\boldsymbol{\beta}^T \cdot \boldsymbol{\beta}} \boldsymbol{\beta} \right| = \frac{2}{|\boldsymbol{\beta}|}$$




Given $\boldsymbol{\beta}$ and b (a boundary), we can:

- Classify all data points (only one boundary is optimal)
- Compute the width of the margin

Now we need to write an optimization problem to search for the optimal $\boldsymbol{\beta}$ and b with the widest margin that classifies all the training data points correctly

Minimize $|\boldsymbol{\beta}|$ Subject to $\boldsymbol{\beta}^T \cdot \mathbf{x} + b > 1$; when in class +1
 $\boldsymbol{\beta}^T \cdot \mathbf{x} + b < -1$; when in class -1

Learning the Maximum Margin Classifier



Minimize $\beta^T \cdot \beta$ Subject to
 $(\beta^T \cdot \mathbf{x}_k + b) y_k \geq 1$ for all (\mathbf{x}_k, y_k)

Incorporating soft constraints

$$\text{Min } \frac{1}{2} \beta^T \cdot \beta + C \sum_{k=1}^R \varepsilon_k \quad \text{s.t.} \quad \begin{aligned} &(\beta^T \cdot \mathbf{x}_k + b) y_k \geq 1 - \varepsilon_k \text{ for all } (\mathbf{x}_k, y_k) \\ &0 \leq \varepsilon_k \leq 1 \end{aligned}$$

Solving the above problem using quadratic programming is a straightforward task


Logistic Regression

- Maximum Likelihood Estimation:

$$\hat{Y} = \arg \max P(Y|X)$$

- We assume that the likelihood function is a logistic function of a linear relationship

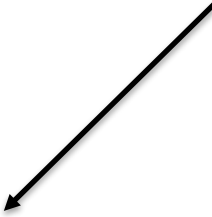
$$P(Y = 1|X) = \frac{1}{1 + e^{-\beta^T X}}$$



Logistic Regression

- Log-likelihood Inference:

Analogous to using Entropy loss


$$\begin{aligned}\hat{\beta} &= \arg \max_{\beta} \log(\Pi_i P(Y_i|X_i)) = \\ &= \arg \max_{\beta} \sum_i \log(P(Y_i|X_i)) = \\ &= \arg \max_{\beta} \sum_i Y_i \log(P(Y_i = 1|X_i)) + (1 - Y_i) \log(P(Y_i = 0|X_i)) = \\ &= \arg \max_{\beta} \sum_i Y_i (-\log(1 + e^{\beta^T X_i})) + (1 - Y_i) (-\log(1 + e^{-\beta^T X_i})) = \\ &= \arg \min_{\beta} \sum_i (1 - Y_i) \beta^T X_i + \log(1 + e^{\beta^T X_i})\end{aligned}$$

- There is no closed form solution:


- We can use numerical optimization, e.g., Newton method
- We can approximate the logit term

Model Evaluation Process

- We use a training dataset for model selection
- A **good** parametric model along with a **suitable** training algorithm guarantees training a model that works well on the training data
- We need to validate that trained models **generalize** well on unseen data instances
- We need a second testing dataset which is fully independent of the training dataset
- We randomly split the annotated dataset into testing and training splits (sometimes, a validation set is generated as well)

Evaluation Metrics

- **Accuracy:** proportion of correctly classified items
 - Accuracy can be dominated by **true negatives** (items correctly classified as not in a class).
 - Sensitive with respect to imbalance
- Precision: $\frac{\text{True Positives}}{\text{True Positives} + \text{False Positive}}$
- Recall: $\frac{\text{True Positives}}{\text{True Positives} + \text{False negative}}$
- Precision and recall are not useful metrics when used in isolation?
- We want our model to have good performance with respect to both metrics
- Implemented in sklearn



Evaluation Metrics

- Why having one measure is helpful?
- $F1 = \frac{2 \text{ Precision Recall}}{\text{Precision} + \text{Recall}}$
- F1 is biased towards the lower of precision and recall:
 - harmonic mean < geometric mean < arithmetic mean
 - F1=0 when Precision=0 or Recall=0
- Generalized F score:

$$F_{\beta} = (1 + \beta^2) \cdot \frac{\text{precision} \cdot \text{recall}}{(\beta^2 \cdot \text{precision}) + \text{recall}}$$