

CSCI 544

Applied Natural Language Processing

Mohammad Rostami
USC Computer Science Department



Logistical Notes

- Project Group Formation Deadline: 09/12 (next week!)
- Do NOT form more than 50 groups
- Meet weekly, helpful for both HW and project
- Paper Selection Deadline: 09/19 (in two weeks)
- Check and then enter your paper:
<https://docs.google.com/spreadsheets/d/1vafG77ijmETCnuVZvKpT35k--5op5wn71GZXgAY700>
- Project Proposal Deadline: 10/03
- HW2:
 - No libraries are allowed, except for common libraries such as Pandas or NumPy

Natural Language Representation

- Language processing hierarchy levels:



- Sparsity in the NLP training datasets: natural language has a very huge space
 - Example: Average Wikipedia page size is 580 words and English has ~1M word roots, yet the actual number of possibilities is far more.
- We need **interpretable** representations or **embeddings** to represent natural language data for model training
 - One-hot representation: too large (15M words) and meaningless
 - Hotel: [0,0,0,0,1,0,0,0,0,0,0,...,0,0,0]
 - Motel:[0,0,0,0,0,0,0,0,0,0,1,0,...,0,0,0]

Similarity of Vectors

- Euclidean distance, i.e., geometric closeness :

- Curse of dimensionality

- Dot product:

$$a \bullet b = ||a|| ||b|| \cos(\theta_{ab})$$

$$= a_1 b_1 + a_2 b_2 + \dots + a_n b_n$$


- Cosine similarity (scale invariant)

$$\cos \theta_{ab} = a \bullet b / ||a|| ||b|| \rightarrow 1 - \cos \theta_{ab} \text{ is a metric}$$

- Invariant with respect to the vector starting point

- **EX:** Hotel: $[0,0,0,0,1,0,0,0,0,0,0,\dots,0,0,0]$, Motel: $[0,0,0,0,0,0,0,0,0,0,1,0,\dots,0,0,0]$

$$\text{Hotel}' * \text{Motel} = 0$$



The Distributional Hypothesis

Zellig Harris,
1954

- Words that occur in the **same contexts** tend to have similar meanings
- Example: nice, good

Budanitsky and
Hirst, 2006

- Word relatedness association: related words **co-occur** in different contexts
- Example: cup, coffee

- If semantic similarity and association of words can be encoded into their representations, we may be able to address the challenge of sparsity
 - In the absence of a particular word during training, we can rely on its synonyms that exist in the training dataset: Motel vs Hotel
 - We can draw conclusions:
Lecturers teach in the university-> Professors ____ in the university.

Vector Embedding of Words

- Represent words using dense vectors:
 - Latent Semantic Analysis/Indexing (SC Deerwester et al, 1988)
 - Word2Vec (Mikolov et al, 2013)
 - GloVe (Pennington et al, 2014)

LSA

- Term weighting-based model
- Consider occurrences of terms at **document level**

Word2Vec

- Prediction-based model
- Consider occurrences of terms at **context level**

GloVe

- Count-based model
- Consider occurrences of terms at **context level**



Word Embedding

- Each word is represented by a vector:
 - The same size is used for all words
 - Relatively low dimensional (~300)
 - Vectors for similar words are similar (measured in dot product)
 - Vector operations can be used for

semantic and syntactic

deductions, e.g.,

Queen – Woman + Man = King



- The key idea is to derive the embeddings from the distributions of word context as they appear in a large corpus.

Singular Value Decomposition

- Every matrix $A \in \mathbb{R}^{m \times n}$ can be factorized as $A = U\Sigma V^T$ where $U \in \mathbb{R}^{m \times m}$ and $V \in \mathbb{R}^{n \times n}$ are orthogonal matrices and $\Sigma \in \mathbb{R}^{m \times n}$ is a diagonal matrix

- The diagonal entries of A are called the singular values of the matrix A

- Singular value $\sigma_i = \sqrt{\lambda_i}$

$$U = AV\Sigma^{-1}$$

$$U\Sigma = AV$$

$$U\Sigma V^T = A$$


Matrix Factorization

- We can form a matrix of M using the idea of Bag of Words: the word representations are highly sparse

	Words											Length of the review(in words)
	1 This	2 movie	3 is	4 very	5 scary	6 and	7 long	8 not	9 slow	10 spooky	11 good	
Review 1	1	1	1	1	1	1	1	0	0	0	0	7
Review 2	1	1	2	0	0	1	1	0	1	0	0	8
Review 3	1	1	1	0	0	0	1	0	0	1	1	6

- Singular value decomposition (U, V are orthonormal)

$$M_{m \times n} = U_{m \times m} \Sigma_{m \times n} V_{n \times n}^T$$



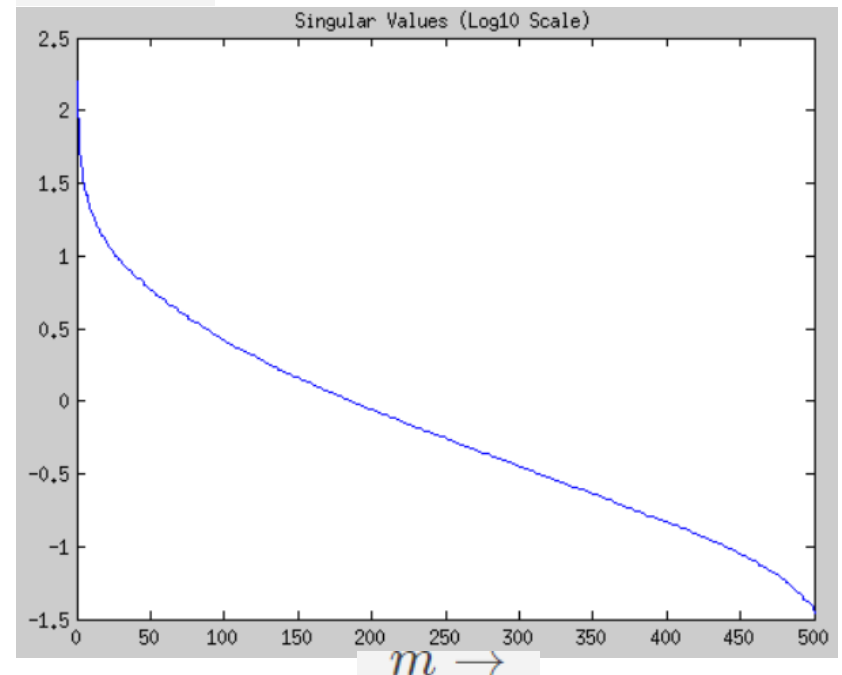
Decrease in σ_m

$$M = U\Sigma V^T$$

$$\Sigma = \begin{bmatrix} \sigma_1 & 0 & 0 & \dots & 0 & \dots & 0 \\ 0 & \sigma_2 & 0 & \dots & 0 & \dots & 0 \\ & & \dots & & & & \\ 0 & 0 & 0 & \dots & \sigma_m & \dots & 0 \end{bmatrix}$$

$$\sigma_1 \geq \sigma_2 \geq \sigma_3 \geq \dots \geq \sigma_m$$

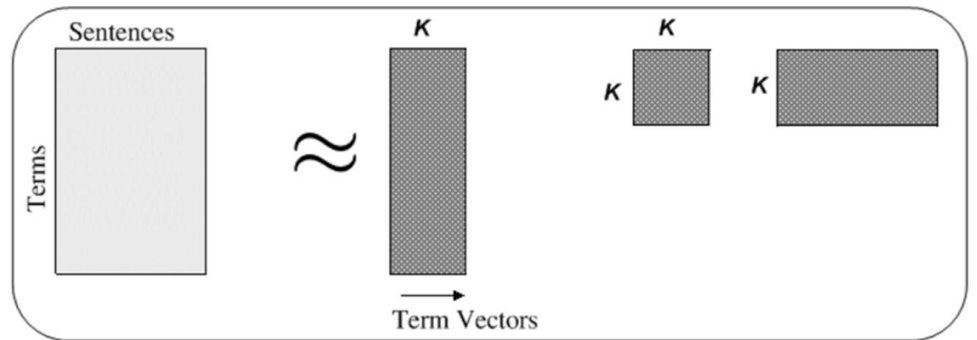
$\log \sigma_m \uparrow$



- What can we do with this information?
 - We can set any sigma after k to be equal to 0
 - Therefore we effectively have a square diagonal matrix of shape $K \times K$ for the new sigma matrix.
 - Because of the 0s in sigma matrix, we can now ignore chunks of U and V matrix as they will result in 0s when we matrix multiply them.

Matrix Factorization

- Many singular values are going to be zero or negligible

$$M_{m \times n} \approx U'_{m \times k} \Sigma'_{k \times k} V'^T_{k \times n}$$


The diagram shows a matrix factorization process. On the left, a large rectangle represents the matrix M with 'Sentences' on the vertical axis and 'Terms' on the horizontal axis. This is followed by an approximation symbol \approx . Then, three smaller rectangles are shown: a tall one labeled K representing U' , a small square labeled K representing Σ' , and a wide one labeled K representing V'^T . Below the tall rectangle is the label 'Term Vectors' with an arrow pointing to it. Below the wide rectangle is the label 'Sentence Vectors' with an arrow pointing to it. To the right of the diagram, two equations are shown:


$$m_l \approx u_l \Sigma' V'^T$$
$$u_l \approx m_l V' \Sigma'^{-1}$$

- We can use rows of U as word embeddings
 - An old idea for dimensionality reduction (it is possible to use other matrix factorization methods, e.g., non-negative matrix factorization)
 - Determining context is heuristic
 - computational expensive with $O(mn^2)$ cost for an $n \times m$ matrix
 - Hard to incorporate new words

Matrix Factorization

- Many singular values are going to be zero or negligible

$$M_{m \times n} \approx U'_{m \times k} \Sigma'_{k \times k} V'^T_{k \times n}$$



- We can use rows of U as word embeddings
- An old idea for dimensionality reduction (it is possible to use other matrix factorization methods, e.g., non-negative matrix factorization)
- Determining context is heuristic
- computational expensive with $O(mn^2)$ cost for an $n \times m$ matrix
- Hard to incorporate new words

(SVD)

Word2Vec

Core idea: find embeddings using a prediction task involving **neighboring words** in a huge real-world corpus.

Concept

Input data: sets of **successive word-patterns** from meaningful sentences in the corpus

We build a **synthetic** prediction task using these patterns

We train a model to solve this prediction task

Embeddings will be the **byproduct** of this task

Example

“One of the most important”

Given an input predict the dash word
Input: [One, of, ___, most, important]
Target: the

$\text{Loss}(\text{Model}([\text{“One”}, \text{“of”}, \text{___}, \text{“most”}, \text{“important”}]), \text{“the”})$

Word2Vec

Concept

Input data: sets of **successive word-patterns** from meaningful sentences in the corpus

We build a **synthetic** prediction task using these patterns

We train a model to solve this prediction task

Embeddings will be the **byproduct** of this task

Specifics

consider a window with the center word w_t and “context words” $w_{t'}$ with a window fixed size, e.g., $(t'=t-5, \dots, t-1, t+1, \dots, t+5)$

predict all $w_{t'}$ given w_t , such that $p(w_{t'} | w_t)$ is maximized

A Two Layer Neural Network

We learn embeddings such that the prediction loss is minimized, i.e., if two words occur in close proximity, their representations become similar.